



Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

LEARNER GUIDE

For

Build a Human-AI Workforce with Autonomous AI Agents

TGS Ref No: TGS-2024043854

Conducted by

Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

Version 1.11

DOCUMENT VERSION CONTROL RECORD

Version Number	Effective Date of Release	Summary of Included Changes	Author
1.0	12 July 2026	Initial release — Learner Guide covering the autonomous-AI-agent labs (Hermes Agent, OpenClaw, Paperclip).	Dr. Alfred Ang
1.1	14 July 2026	Revised for the 2-day structure and 34 labs (Hermes 14, OpenClaw 12, Paperclip 8) with per-lab reference videos; learning outcomes aligned to the Course Proposal / TSC.	Dr. Alfred Ang
1.2	14 July 2026	Per-lab step numbering; YouTube links removed (videos remain in the labs); Hermes install lab expanded with the official install methods.	Dr. Alfred Ang
1.3	14 July 2026	Deck restyled in the dark masterclass theme end-to-end; added the Hermes masterclass slides (What is Hermes; The Bet: Compounding Value; platform matrix; Top-10 slash commands; Telegram setup; memory L1-vs-L2 and session-search; what-is-a-tool, built-in tool catalog, tools-in-action, cron scheduling, delegation, kanban orchestration and OpenClaw topic slides); Lab 13 rebuilt as the multi-agent video team (profiles + kanban) with detailed Labs 11-14 guides; fade slide transitions and the live Achievements dashboard screenshot.	Dr. Alfred Ang
1.4	15 July 2026	Security masterclass slides added to the Hermes Security lab (defense-in-depth layers; the gateway's six-check default-deny	Dr. Alfred Ang

		trust chain; channel lockdown dials; the three approval modes + the Tirth scanner; a live Dangerous-Command approval; YOLO mode; the unrecoverable-command blocklist floor; prompt-injection scanning of context files). Security lab steps expanded with the channel allowlist and approvals.mode.	
1.5	15 July 2026	Sixteen concept diagrams imported from the original Mastering OpenClaw masterclass deck into Topic 2 (OpenClaw OS control plane; deployment; hub-and-spoke gateway; tools vs skills; SKILL-on-demand; 6-phase execution loop; exec; HEARTBEAT + autonomic tick; workspace file persistence; context assembly; compaction; defense methods; poisoned-skill anatomy; sessions_spawn; sub-agent routing).	Dr. Alfred Ang
1.6	15 July 2026	Topic 3 (Paperclip) rebuilt as the 'Tertiary AI News Research' company across 6 labs (27-32): install on Windows & Mac; company, CEO & mission; model adaptors; a detailed task backlog whose core task hires the team; the Tavily Search API via Secrets; and hiring the members under the hiring task — with live dashboard screenshots. Natural-language prompt steps added to the Hermes and OpenClaw labs alongside the CLI.	Dr. Alfred Ang

		Platform logos added to the Three Platforms page. 32 labs total.	
1.7	15 July 2026	Sixty-eight content slides migrated from the original 'Mastering OpenClaw' masterclass deck into Topic 2 as framed figures in their matching labs — agent evolution and Agent = Model + Harness, the OpenClaw Wikipedia history and OpenAI acquisition, deployment options and services, model stacks and Mission Control, channels, skills and ClawHub risks, tools (AgentMail, Agent Browser, Firecrawl), system prompt / workspace / context and memory, security lockdown, slash commands, crons, and use cases (NFT certificates, dashboard, cost saving, Google integration).	Dr. Alfred Ang
1.8	15 July 2026	Imported masterclass figures re-blended into the dark deck theme: original title and copyright strips cropped, white canvases converted to the deck background with text inverted to ivory (embedded screenshots untouched), frameless rendering, and duplicated titles reworded.	Dr. Alfred Ang
1.9	16 July 2026	Deck curated per trainer review: the bulk masterclass figure imports were removed, keeping nine key figures (Wikipedia history, Agent=Model+Harness, SKILL.md, ClawHub, AgentMail, 10-steps lockdown,	Dr. Alfred Ang

		HEARTBEAT, autonomic tick, workspace persistence) re-exported at high resolution, with the Topic 2 opening story reordered ahead of Key Concepts. Every CLI step now also shows the natural-language prompt alternative ('Or just ask') across the Hermes and OpenClaw labs, in the deck and Learner Guide.	
1.10	16 July 2026	PP assessment updated to v4 — evidence screenshots changed to: Task 1 Hermes interacting on Telegram or the multi-agent Kanban board; Task 2 OpenClaw on Telegram or the TUI; Task 3 the Paperclip setup (company, mission, CEO, hired agents) — with Topic 3 lab references refreshed. A 'RAG vs a Standard LLM Assistant' concept slide added to Core Concepts to ground WA Q2. Final PPT/LP/LG/labs/assessment alignment review.	Dr. Alfred Ang
1.11	16 July 2026	Topic 3 company renamed to 'Alterra AI Blogs' with the mission re-pointed at publishing AI-related blogs end to end; the core backlog task is now 'Hire the agents to publish the AI-related blogs end to end' and Lab 32's hiring happens explicitly under it; Lab 29 (Adaptor) and Lab 32 descriptions expanded in full; NEW Lab 33 'Routine & Trigger — Publish Daily at 3pm'	Dr. Alfred Ang

		(33 labs total), verified with a live 'Run now' pipeline test and two live screenshots.	
--	--	---	--

TABLE OF CONTENTS

Introduction.....	8
Course Learning Outcomes.....	8
Before You Start — Environment Setup.....	8
Topic 01 — Hermes Agent — Your Autonomous Personal Chief-of-Staff (33%).....	9
Lab 1 — Install and Setup Hermes.....	10
Lab 2 — Deployment — Local Install & Hermes Desktop App.....	11
Lab 3 — Memory and Plugins.....	12
Lab 4 — Skills.....	13
Lab 5 — Providers & Model.....	14
Lab 6 — MCP and Tools.....	15
Lab 7 — Cron Jobs & Automation.....	16
Lab 8 — Subagents & Delegation.....	17
Lab 9 — Profile and Kanban.....	18
Lab 10 — Security.....	18
Lab 11 — Use Case — Make a Video with Hyperframe.....	20
Lab 12 — Visualization.....	21
Lab 13 — Build a Multi-Agent Video Team (Profiles + Kanban).....	23
Lab 14 — Webhook.....	25
Topic 02 — OpenClaw — Automating a Business Back-Office (33%).....	27
Lab 15 — Install OpenClaw.....	27
Lab 16 — Models & Providers.....	29
Lab 17 — Channels.....	31
Lab 18 — Skills.....	32
Lab 19 — Tools & Integrations.....	34
Lab 20 — Commands.....	35
Lab 21 — Cron Jobs & Heartbeat.....	37
Lab 22 — Memory & Context.....	40
Lab 23 — Security.....	41
Lab 24 — Multi-Agent Workforce.....	43
Lab 25 — Dreaming — Idle Reflection & Self-Improvement.....	46
Lab 26 — Use Cases & Key Functions.....	47
Topic 03 — Paperclip — Running & Governing a Company of AI Agents (34%).....	48
Lab 27 — Install Paperclip on Windows & Mac.....	48
Lab 28 — Setup Company, CEO & Mission.....	49
Lab 29 — Setup Adaptor.....	50
Lab 30 — Create the Task Backlog.....	51
Lab 31 — Add the Tavily Search API.....	52

Lab 32 — Hire the Members Under the Hiring Task.....	53
Lab 33 — Routine & Trigger — Publish Daily at 3pm.....	55
Wrap-Up — The Agent Build Arc and Governance.....	56
Next Steps.....	57
Glossary.....	57

Introduction

This Learner Guide accompanies the WSQ course Build a Human-AI Workforce with Autonomous AI Agents (TGS-2024043854), conducted by Tertiary Infotech Academy Pte Ltd. It provides step-by-step instructions for all 33 hands-on labs across three autonomous-AI-agent platforms — Hermes Agent, OpenClaw and Paperclip — organised into three topics. Each platform builds on the same skillset: install the runtime, connect models and channels, give the agent memory, extend it with skills and tools, automate it with schedules, secure it with governance, and finally orchestrate multiple agents to deliver a realistic business outcome.

Use this guide alongside the course slides and the lab files in the labs/ folder of the course repository. The labs run on your own laptop and in Docker; where a lab connects to an external service (an LLM provider, a messaging channel, a VPS) use only accounts and credentials you own, and keep API keys and tokens out of prompts and out of version control.

Course Learning Outcomes

- LO1: Analyse AI (LLM-based) applications across a range of industries to identify their capabilities and limitations.
- LO2: Establish the relationship between AI algorithm/agent design and chatbot/agent efficiency.
- LO3: Evaluate and improve the effectiveness of AI (RAG and agent) applications in product development.

Before You Start — Environment Setup

What you need

- A laptop you have administrator rights on — Windows 10/11, macOS 12+ or Ubuntu 22.04+ — with a terminal.
- Node.js 24 LTS (required to run the OpenClaw gateway daemon).
- Docker Desktop (required to self-host Paperclip with Docker Compose, and used for isolated agent sandboxes).
- At least one LLM provider login or API key — Claude (Anthropic), OpenAI, OpenRouter, MiniMax or DeepSeek.
- A free Telegram account for the messaging-channel labs (you will create a bot via BotFather).
- Optional: a VPS (e.g. Hostinger or exe.dev) if you want to run an agent 24/7 beyond the classroom.

Verify your toolchain

Open a terminal and confirm the core prerequisites are present before you begin. Each platform's installer (Hermes, OpenClaw, Paperclip) is run in its own lab; here you only confirm the building blocks.

```
$ node --version      # Node.js 24.x for OpenClaw
$ docker --version    # Docker Desktop for Paperclip and sandboxes
$ docker compose version
$ git --version        # to clone the course labs repository
```

Conventions used in every lab

- Commands are run from your terminal; a leading `sudo` is used only where elevated rights are required.
- Placeholders such as `<API_KEY>`, `<BOT_TOKEN>` and `<HOST>` are replaced with your own values.
- Store provider keys and channel tokens in each platform's config (never in prompts or in git).
- Keep approval prompts, budgets and sandboxing on so agents act under human oversight.
- Each capstone orchestrates multiple agents — start the single-agent labs first so the pieces are familiar.

Topic 01 — Hermes Agent — Your Autonomous Personal Chief-of-Staff (33%)

Install & Setup · Deployment · Memory & Plugins · Skills · Providers & Model · MCP & Tools · Crons · Subagents · Profile & Kanban · Security

Key concepts

- Hermes Agent (Nous Research) installs on your own machine — via the install script or the Hermes Desktop app — and requires a model with a $\geq 64,000$ -token context window.
- Memory and plugins make the agent stateful and extensible: a three-layer memory (agent-curated notes, FTS5 cross-session recall, Honcho user modelling) plus installable skills and MCP tools.
- Providers & models are swappable with no lock-in (Nous Portal, OpenRouter, OpenAI, any endpoint); MCP servers and the Tool Gateway (web search, image, TTS) give the agent real-world reach.
- Cron jobs automate recurring work; subagents & delegation split work across isolated backends; the profile and Kanban board let you supervise the agent's tasks — all under security controls.

Lab 1 — Install and Setup Hermes

Learning outcome: LO1: Install and configure Hermes Agent on a local machine..

Goal: Install Hermes Agent on your laptop using any of the official install methods (Desktop installer, install script, PowerShell, from source, or Nix), run the first-time setup wizard, and confirm a healthy install. This is the foundation for the Athena chief-of-staff agent you build across the track. Prerequisite: Git; the installer auto-handles Python 3.11, Node.js v22, ripgrep and ffmpeg.

What you'll build

A working Hermes Agent install that passes 'hermes doctor' and opens the TUI. (Tools: Hermes Agent, Nous Portal, terminal.)

Step-by-step

1. Choose your installation method Hermes offers several official install paths — pick the ONE that matches your machine: > If you install via the Desktop app (Option A), you can skip Steps 2–3. If you install via the CLI first, you can add the desktop app later with hermes desktop.
2. Option B — install script (Linux / macOS / WSL2 / Termux) Run the official one-line installer. It downloads the Hermes runtime, places it under your user directory, and wires up the hermes command:

```
curl -fsSL https://hermes-agent.nousresearch.com/install.sh | bash
```

3. Option C — Windows PowerShell (native Windows) On native Windows (no WSL2), run the PowerShell installer instead:

```
iex (irm https://hermes-agent.nousresearch.com/install.ps1)
```

4. Install Hermes Desktop (macOS / Windows) Hermes Desktop gives you a GUI over the same agent. Either download the installer from <<https://hermes-agent.nousresearch.com/>> and run it, or — if you already installed the CLI in Step 2/3 — launch it directly:

```
hermes desktop
```

5. Reload your shell so the hermes command is on PATH The installer adds Hermes to your PATH via your shell profile, but your current terminal session hasn't re-read that file yet. Reload it:

```
source ~/.zshrc
```

6. Run the first-time setup wizard Launch the interactive setup wizard and connect Hermes to your Nous Portal account. This is where you sign in, pick defaults, and let Hermes provision a model provider:

```
hermes setup --portal
```

7. Check the install is healthy Run the built-in diagnostic. It checks the binary, config, PATH, provider connectivity, and the local runtime:

```
hermes doctor
```

8. Launch the agent and say hello Open the terminal UI (TUI) and start your first conversation with the agent — this is Athena's very first session:

```
hermes --tui
```

Test it

'hermes doctor' reports all green and the TUI opens and replies to a message.

Note: Full commands and screenshots are in labs/lab01-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 2 — Deployment — Local Install & Hermes Desktop App

Learning outcome: LO1: Deploy Hermes locally and run it through the Hermes Desktop app..

Goal: Run Hermes as a persistent local deployment (preferred) and also install the Hermes Desktop app for a GUI experience. Confirm the local gateway is serving and keep the runtime up to date.

What you'll build

A local Hermes deployment plus the Desktop app, both talking to the same agent.
(Tools: Hermes Agent, Hermes Desktop, terminal.)

Step-by-step

1. Watch the reference videos Watch both videos above. The primary shows the local deployment; the additional walkthrough covers the Desktop app in more depth. Note any UI screens that differ from the steps below.
2. Confirm the local install runs (preferred deployment) Verify the CLI runtime you installed in Lab 1 is present and check its version. The local install is the preferred deployment — it keeps the agent runtime on your machine:

```
hermes --version
```

3. Download and install the Hermes Desktop app Download the Desktop app for your OS from the official site and install it like any normal application:
4. Open the Desktop app and sign in Launch the Hermes Desktop app and sign in with the same Nous Portal account you used for the CLI in Lab 1. Signing in with the same account is what makes the Desktop app and the CLI drive the same agent. > This step is UI-only — there is no terminal command. In the app, choose Sign in with Nous Portal and complete the browser/authentication prompt.

5. Verify the local gateway is serving The gateway is the local service that both the CLI and Desktop app connect to. Confirm it is up and healthy: Or just ask the agent: "Check that your local gateway is up and serving"

`hermes gateway status`

6. Keep the runtime up to date Update Hermes to the latest release so the CLI and Desktop app stay compatible:

`hermes update`

Test it

Both the local CLI and the Desktop app drive the same agent; 'hermes gateway status' is healthy.

Note: Full commands and screenshots are in labs/lab02-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 3 — Memory and Plugins

Learning outcome: LO2: Give the agent persistent memory and extend it with plugins..

Goal: Explore Hermes' three-layer memory (agent-curated notes, FTS5 cross-session recall, Honcho user modelling). Teach Athena a preference, recall it in a fresh session, and enable a plugin to extend behaviour.

What you'll build

Athena remembers a stated preference across sessions and runs an enabled plugin. (Tools: Hermes Agent, memory, plugins.)

Step-by-step

1. Teach Athena a durable preference Start a session (`hermes --tui`) and give Athena a clear, memorable instruction that should persist. Explicitly asking it to remember nudges the agent to write the fact into its curated-notes memory layer:

`Remember I prefer concise, bulleted summaries.`

2. Start a new session and confirm the preference is recalled Resume the conversation in a fresh session. The `--continue` flag brings back prior context so you can test cross-session recall: Or just ask the agent: "What summary style did I ask you to remember?"

`hermes --continue`

3. Inspect stored memory / sessions List the stored sessions so you can see the memory Hermes is persisting behind the scenes: Or just ask the agent: "List my recent sessions and what each was about"

`hermes sessions list`

4. Enable a plugin to extend the agent Enable a plugin to add capability beyond the base agent. Plugins are toggled from the Hermes plugins UI or config:

Test it

A preference taught in one session is recalled in a new session, and the enabled plugin is active.

Note: Full commands and screenshots are in `labs/lab03-*/README.md`. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 4 — Skills

Learning outcome: LO3: Extend the agent with installable, self-improving skills..

Goal: Browse and install open-standard skills (agentskills.io / Skills Hub) so Athena can perform new tasks. Skills are portable and can self-improve as the agent uses them.

What you'll build

Athena equipped with at least one installed skill that it can invoke on request. (Tools: Hermes Agent, agentskills.io, Skills Hub.)

Step-by-step

1. Browse the skills hub Open the skills catalogue to see what's available to install: Or just ask the agent: "Show me what skills are available on the skills hub"

`hermes skills browse`

2. Search for a skill by keyword Narrow the catalogue to something useful for a chief-of-staff agent — for example calendar, email, research, or pdf: Or just ask the agent: "Search the skills hub for <keyword> skills"

`hermes skills search <keyword>`

3. Install a skill Install a skill by its full identifier (copy it from the search results). The pattern is `owner/skills/name`: Or just ask the agent: "Install the <name> skill from the skills hub and confirm it loaded"

`hermes skills install <owner/skills/name>`

4. Ask the agent to use the new skill and observe self-improvement Open the TUI and ask Athena to perform a task that the new skill enables. Watch it select and run the skill: > This step is conversational — no fixed command. Phrase a request that maps to the skill (e.g. for a summarize skill: "Summarize this article for me: <paste text>"). As the agent uses the skill repeatedly, it refines how it invokes it — that's the self-improvement aspect.

5. The natural-language way You can skip the CLI entirely and simply ask the agent: "Learn the slide-deck skill from agentskills.io and confirm it is installed".

Test it

The installed skill appears in the agent's toolset and runs when invoked.

Note: Full commands and screenshots are in labs/lab04-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 5 — Providers & Model

Learning outcome: LO2: Connect an LLM provider and switch models with no lock-in..

Goal: Connect a model provider (Nous Portal, OpenRouter, OpenAI or any endpoint) and select a model that meets the $\geq 64,000$ -token context requirement. Switch providers to compare speed and quality.

What you'll build

Athena running on a chosen provider/model, with the ability to switch on demand.
(Tools: Hermes Agent, Nous Portal, OpenRouter.)

Step-by-step

1. Open the interactive provider/model picker Launch the interactive picker to see available providers and models and select one: Or just ask the agent: "Show me the available providers and models, and which one you're on"

```
hermes model
```

2. Set a model explicitly Instead of (or in addition to) the picker, set the active model directly in config: Or just ask the agent: "Switch yourself to Claude Opus 4.6"

```
hermes config set model anthropic/claude-opus-4.6
```

3. Add a provider key if using a cloud endpoint If your chosen model runs on a cloud endpoint (e.g. OpenRouter), store the API key in config so Hermes can authenticate: Or just ask the agent: "Add my OpenRouter API key to your provider config"

```
hermes config set OPENROUTER_API_KEY sk-or-...
```

4. Confirm the active model meets the $\geq 64k$ context rule Reopen the picker to confirm the active model and its context window: Or just ask the agent: "Confirm your active model has at least a 64k context window"

```
hermes model
```

Test it

The agent starts on the selected model and can be switched to another provider without reinstalling.

Note: Full commands and screenshots are in labs/lab05-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 6 — MCP and Tools

Learning outcome: LO3: Give the agent real-world reach with MCP servers and the Tool Gateway..

Goal: Add Model Context Protocol (MCP) servers in the Hermes config and use the bundled Tool Gateway (web search, image generation, TTS) so Athena can act on the outside world.

What you'll build

Athena wired to an MCP server (e.g. GitHub) and able to use a Tool Gateway tool. (Tools: Hermes Agent, MCP, Tool Gateway.)

Step-by-step

1. Open the Hermes config to add an MCP server Edit your Hermes config file and locate (or create) the `mcp_servers:` section:
2. Add a GitHub MCP server entry Add an entry that launches the official GitHub MCP server via `npx`. A typical block looks like this: Or just ask the agent: "Add the GitHub MCP server to your tool config"

`mcp_servers:`

`github:`

`command: npx`

`args: ["-y", "@modelcontextprotocol/server-github"]`

`env:`

`GITHUB_PERSONAL_ACCESS_TOKEN: "ghp_your_token_here"`

3. Restart so the MCP server is picked up, then verify Restart Hermes (close and reopen the TUI, or restart the gateway) so it reads the new config, then run the diagnostic: Or just ask the agent: "Run a health check and confirm the GitHub MCP tools are loaded"

`hermes doctor`

4. Use a Tool Gateway tool from a chat Open the TUI and ask Athena to use a bundled Tool Gateway capability — web search, image generation, or text-to-speech: > This step is conversational. For example: "Search the web for today's top AI agent news and list 3 headlines," or "Generate an image of a friendly robot assistant." Confirm the agent invokes the tool and returns the result.

5. The natural-language way You can skip the CLI entirely and simply ask the agent: "Add the GitHub MCP server to your tools and list what it can do".

Test it

The agent lists the new MCP server's tools and completes a task using a Tool Gateway tool.

Note: Full commands and screenshots are in labs/lab06-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 7 — Cron Jobs & Automation

Learning outcome: LO4: Automate recurring agent work on a schedule..

Goal: Schedule Athena to run recurring jobs — for example an 8am daily briefing sent to your messaging app — and verify the scheduled run fires and produces output.

What you'll build

A scheduled job (e.g. a daily briefing) that runs automatically and delivers a result. (Tools: Hermes Agent, scheduler/crons.)

Step-by-step

1. Define a recurring job (daily 8am briefing) Create a recurring automation in the Hermes automations/crons section — for example a daily 8am briefing:
2. Confirm the automation is registered Run the diagnostic to confirm Hermes has picked up the new scheduled job: Or just ask the agent: "Check that my daily 8am briefing automation is registered"

`hermes doctor`

3. Trigger the job once to confirm it produces output Manually run the job once (rather than waiting for 8am) to confirm it produces the expected briefing. > This is typically a "run now" action in the automations UI, or a trigger command shown in the docs. Confirm the briefing is generated and delivered to your chosen destination.
4. Confirm the next scheduled run is queued Check that after the manual run, the next run is still scheduled for its normal time (08:00 tomorrow), so the automation keeps recurring.
5. The natural-language way You can skip the CLI entirely and simply ask the agent: "Every weekday at 8am, summarize the top AI news and send it to me on Telegram" — the agent schedules itself with its cronjob tool.

Test it

The scheduled job runs at its time (or on manual trigger) and delivers the expected briefing.

Note: Full commands and screenshots are in labs/lab07-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 8 — Subagents & Delegation

Learning outcome: LO5: Delegate work to isolated subagents across backends..

Goal: Turn Athena into a coordinator that delegates tasks to worker subagents running on isolated terminal backends (local, docker, ssh, modal), then aggregates their results.

What you'll build

A coordinator agent that delegates a task to one or more worker subagents and combines the output. (Tools: Hermes Agent, terminal backends (docker/ssh/modal).)

Step-by-step

1. Choose an isolated backend for worker agents Set the terminal backend that worker subagents will run on. docker gives each worker an isolated container: Or just ask the agent: "Run your worker subagents in the docker sandbox from now on"

```
hermes config set terminal.backend docker
```

2. Ask the coordinator to delegate a multi-part task Open the TUI and give Athena (the coordinator) a task with clearly separable parts so it delegates each to a worker subagent: > Conversational step. For example: "Delegate this: subagent A researches competitor pricing, subagent B drafts a comparison table, subagent C writes a one-paragraph recommendation. Then combine their work." Watch the coordinator spin up workers.
3. Observe each subagent run in isolation and report back Watch each worker subagent execute on the isolated backend and return its portion of the result. Confirm they run separately (each in its own container/environment when using docker) rather than in the coordinator's session.
4. Review the coordinator's aggregated result Confirm the coordinator combines the workers' outputs into one coherent deliverable — the research, the table, and the recommendation merged into a single response.
5. The natural-language way You can skip the CLI entirely and simply ask the agent: "Delegate three subagents to research MiniMax, Kimi and DeepSeek in parallel, then merge the findings".

Test it

The coordinator delegates to worker subagents on an isolated backend and returns a combined result.

Note: Full commands and screenshots are in labs/lab08-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 9 — Profile and Kanban

Learning outcome: LO5: Supervise the agent through its profile and Kanban board..

Goal: Configure the agent's profile and use the Hermes Kanban board to watch tasks move through todo -> in progress -> done, giving you human oversight of what Athena is working on.

What you'll build

A configured agent profile and a Kanban board tracking the agent's live tasks. (Tools: Hermes Agent, profile, Kanban board.)

Step-by-step

1. Set up the agent profile (identity, defaults, preferences) Configure Athena's profile — its identity, default model/backend, and behavioural preferences. > Open the profile settings in the Desktop app (or the profile section of `~/.hermes/config.yaml`). Set the name to Athena, role to personal chief of staff, and default preferences (concise bulleted output, timezone, working hours). The exact profile fields are version-specific — verify in the video / Hermes docs (<https://hermes-agent.nousresearch.com/docs/>).
2. Open the Kanban board view Open the Kanban board that visualizes the agent's tasks. > In the Desktop app, open the Kanban / Tasks view. You should see columns such as Todo, In Progress, and Done.
3. Assign a task and watch it move todo → in progress → done Give Athena a task and watch the corresponding card move across the board: > Assign a task (e.g. "Draft my weekly team update"). A card appears in Todo, moves to In Progress as the agent works, and lands in Done when finished. This is your live oversight of what the agent is doing.
4. Use the board to review and accept completed work When a card reaches Done, open it, review the agent's output, and accept (or send it back with feedback). This human-in-the-loop review is the point of the board.

Test it

The profile is set and the Kanban board shows tasks progressing across its columns.

Note: Full commands and screenshots are in `labs/lab09-*/README.md`. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 10 — Security

Learning outcome: LO4: Secure the agent with isolation, secrets and approvals..

Goal: Harden Athena with defense in depth: restrict who can message the agent (default deny), run risky work in an isolated terminal backend, manage secrets safely, and require approval before sensitive actions — applying least privilege throughout.

What you'll build

A hardened agent that allowlists its users, sandboxes execution, protects secrets and prompts for approval on risky actions. (Tools: Hermes Agent, gateway allowlists, terminal backends, secrets management.)

Step-by-step

1. Restrict who can message the agent — set the channel allowlist (default deny)
Decide who is allowed to talk to Athena. The gateway authorizes every inbound message with a six-check chain — per-platform allow-all, DM pairing, the platform allowlist, the global allowlist, global allow-all, then default deny. Nothing configured means everyone is denied (fails closed), with a startup warning telling you exactly that. In the dashboard open Channels → Configure for your channel (e.g. Telegram) and set the allowlist to your own user ID: - Allowed user IDs (TELEGRAM_ALLOWED_USERS) — comma-separated IDs that may use the bot (get yours from @userinfobot). - Allow all users (TELEGRAM_ALLOW_ALL_USERS) — leave off; it opens the bot to anyone (dev only). - Unknown DMs can also be admitted one-by-one via pairing codes you approve.
2. Isolate execution in a sandboxed backend Set the terminal backend to a sandboxed container so risky commands never run directly on your host: Or just ask the agent: "Run your risky work in the docker sandbox from now on"

```
hermes config set terminal.backend docker
```

3. Store provider/tool secrets via config rather than plain text Store API keys through the Hermes config mechanism instead of pasting them into prompts or scripts: Or just ask the agent: "Store this <PROVIDER> API key securely in your config"

```
hermes config set <PROVIDER>_API_KEY <value>
```

4. Require approval before the agent runs risky actions — set the approval mode Set the approval mode so the agent pauses and asks you before executing sensitive actions (shell commands, file deletions, external posts): Or just ask the agent: "Ask me for approval before running any risky action"

```
hermes config set approvals.mode manual
```

5. Apply least privilege — grant only the tools/skills the task needs Review the tools, skills, and MCP servers Athena has access to and disable anything not needed for the current work. > Conversational/config step: trim the enabled skills (Lab 4) and MCP servers (Lab 6) to the minimum. Least privilege means the agent can only do what the task actually requires.

Test it

Only allowlisted users reach the agent; risky actions run in the sandbox and only after approval; secrets are not stored in plain text.

Note: Full commands and screenshots are in labs/lab10-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 11 — Use Case — Make a Video with Hyperframe

Learning outcome: LO5: Apply the agent to a real creative use case — produce a video with Hyperframe..

Goal: Put Athena to work on an end-to-end creative task: use the agent to drive the Hyperframe tool and generate a short video from a concept brief, then review and refine the result.

What you'll build

A short video generated end-to-end by the agent using Hyperframe. (Tools: Hermes Agent, Hyperframe.)

Step-by-step

1. Install / verify the Hyperframe video skill Wire Hyperframe into Athena the same way you added skills in Lab 4. In a terminal, search the skills hub for the Hyperframe (video) skill and install the match; if your build integrates Hyperframe as an MCP server or Tool Gateway tool instead, add it the way you added tools in Lab 6. Or just ask the agent: "Find the Hyperframe video skill on the skills hub and install it"

```
hermes skills search hyperframe
```

```
hermes skills install <owner/skills/hyperframe>
```

2. Set the Hyperframe credential via config If Hyperframe needs an API key or account token, store it through Hermes config so it never sits in plain text. Get the key from your Hyperframe account page, then set it and confirm no error is printed. Or just ask the agent: "Save my Hyperframe API key securely in your config"

```
hermes config set HYPERFRAME_API_KEY <your-key>
```

3. Confirm the tool is loaded, then start a session Run the health check and confirm it reports green with the new skill/tool listed. Then open a chat session (TUI or Desktop app) and ask Athena "What video tools do you have?" — it should name Hyperframe in its reply. If it doesn't, restart Hermes so the new tool is picked up. Or just ask the agent: "Run a health check and confirm the Hyperframe skill is loaded"

```
hermes doctor
```

```
hermes --tui
```

4. Brief the agent with a concrete 3-sentence brief Type a tight, three-sentence creative brief covering topic, style/tone, and length + extras. A concrete brief is the single biggest lever on output quality — vague briefs produce generic videos. For example: > "Make a 30-second explainer video introducing 'Athena, my AI chief of staff'. Clean, modern visual style with an upbeat tone. Keep it to 30 seconds and add on-screen captions."
5. Have the agent generate the video and wait for the render Ask Athena to generate the video now. You should see the agent invoke the Hyperframe tool in the session transcript (a tool-call entry), then wait — rendering typically takes a few minutes for a 30-second clip. When it finishes, Athena reports back with a link or file path to the rendered video; if the tool call errors, check the credential from step 3.
6. Play the result, give one round of feedback, and regenerate Open and play the video end-to-end. Then give Athena one specific round of feedback — change exactly the things that bothered you, e.g. "Slow the pacing, make the captions larger, and warm up the colour palette" — and ask it to regenerate. Compare the two versions; you should see your feedback reflected in version 2.
7. Export / save the final video and note where the file lands Ask Athena to save/export the final cut and tell you the exact file path. Outputs typically land in the session's workspace under your Hermes home directory — open it and confirm the file plays from disk: Or just ask the agent: "Where did you save the final video? Give me the file path"

```
open ~/.hermes/
```

Test it

The agent produces a playable video from your brief using Hyperframe, and you can locate the saved file.

Note: Full commands and screenshots are in labs/lab11-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 12 — Visualization

Learning outcome: LO5: Visualize agent activity and workflow data..

Goal: Use Hermes' visualization capability to see the agent's workflow, task flow and data at a glance, making the agent's behaviour transparent and easier to supervise.

What you'll build

A visualization of the agent's workflow / activity you can read and interpret. (Tools: Hermes Agent, visualization.)

Step-by-step

1. Run a real multi-step session so there is something to visualize A visualization of an empty agent is empty. Start a fresh session and give Athena a task that forces several tool calls in sequence — for example: > "Research the top 3 AI agent platforms, compare their pricing, and summarize in a table." Let it run to completion; the web searches, reads, and summarization steps it performs become the nodes of your graph.
2. List sessions and note the ID of the one to visualize In a terminal, list your recent sessions and copy the ID (or title/timestamp) of the session you just ran. You will visualize this specific session rather than "everything", which keeps the diagram readable. Or just ask the agent: "List my recent sessions so I can pick one to visualize"

`hermes sessions list`

3. Open the visualization view In the Desktop app, open the Visualization / Activity view from the sidebar and select the session from step 3. You should see the session's activity surface — task flow, tool calls, and timing. > The exact menu location and view name are version-specific — verify in the video / Hermes docs (<https://hermes-agent.nousresearch.com/docs/>).
4. Generate a workflow graph of the session Ask Athena (or use the view's generate button) to render the session as a workflow diagram: > "Visualize the steps you took in this session as a workflow diagram." A graph should render in which each node is a tool call (web search, file read, summarize...) and edges show the order the calls ran in. If nothing renders, confirm you selected a session that actually used tools.
5. Read the graph — nodes, sequence, and time Walk the diagram from start to finish and name what each node did: which tool, with what input, producing what output. Check the sequence matches what you asked for, and look for where time was spent (long-running nodes) or where the agent looped/retried. This is the supervision skill: confirming from the graph that the agent did what you expected — nothing more, nothing less.
6. Export or screenshot the visualization Save the diagram for your records — use the view's export button if your build has one, otherwise take a screenshot. Keep it with the lab evidence; you will reuse this technique whenever you need to audit an agent run.
7. Use the graph to explain what the agent did Turn to a classmate (or write three sentences) and explain the run using only the graph: what the goal was, which tools ran in what order, and where the result came from. If you can narrate the run from the diagram alone, the visualization has done its job of making Athena's behaviour transparent.

Test it

A workflow visualization of a real session renders; you can name each node (tool call) and explain what the agent did.

Note: Full commands and screenshots are in labs/lab12-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 13 — Build a Multi-Agent Video Team (Profiles + Kanban)

Learning outcome: LO3: Orchestrate a team of profiles through the Kanban board to produce a video..

Goal: Build a team of three Hermes profiles — a researcher (long-context model) that reads sources and writes the video brief, a scriptwriter (cheap model) that turns the brief into a narrated script, and a video_producer (code/tool model) that renders the video with TTS and slides — then create a Kanban board task, decompose it into child tasks, and let the orchestrator route each child task to the right profile by its description. The board is a durable SQLite database (~/.hermes/kanban.db) and each task gets its own workspace.

What you'll build

A three-profile video team whose Kanban tasks flow triage -> todo -> ready -> running -> done and deliver a rendered video. (Tools: Hermes Agent, profiles, Kanban board, dashboard.)

Step-by-step

1. Design the team on paper first Before creating anything, write down the three roles, the kind of model each needs, and — most importantly — the one-line description the orchestrator will route by. The description is the routing contract: it must say exactly what the profile takes in and hands off. Note the pipeline shape: brief → script → video. Each profile's output is the next profile's input.
2. Create the researcher profile Create the first profile on a long-context model (it will read sources whole). The --description text is what the orchestrator matches child tasks against, so type it exactly as designed: Or just ask the agent: "Create a researcher profile on a long-context model that writes video briefs"

```
hermes profile create researcher --description "Reads sources, writes the video brief + key points"
```

3. Create the scriptwriter profile Create the second profile on a cheap model — turning a brief into a script is high-volume, low-difficulty work, so this is where you save money: Or just ask the agent: "Create scriptwriter and video_producer profiles for the script and the render"

```
hermes profile create scriptwriter --description "Turns the brief into a narrated script + scene list"
```

4. Create the video_producer profile Create the third profile on a code/tool-capable model — it must drive TTS and slide-rendering tools reliably:


```
hermes profile create video_producer --description "Turns the script into a rendered video with TTS + slides"
```

5. Verify the team on the Profiles page Open the dashboard and go to the /profiles page in the sidebar. You should see all three profiles listed, each showing its model and its description. Read each description back critically — a vague description here causes wrong routing later. Fix any typos now (edit the profile or recreate it).

6. Create the Kanban board task Create the top-level task on the board and assign the first stage to the researcher. Pick a real topic you care about and keep the deliverable measurable ("60-second explainer video"): Or just ask the agent: "Add a Kanban task: produce a 60-second explainer on <topic>, assign it to the researcher"

```
hermes kanban create "Produce a 60-second explainer video on <topic>" --assignee researcher
```

7. Decompose the task into child tasks Ask Hermes to break the top-level task into child tasks: Or just ask the agent: "Decompose that board task into child tasks and route each to the right profile"

```
hermes kanban decompose <id>
```

8. Watch the lanes as the team works Open the dashboard /kanban page and watch the board dispatch. Every task moves through the lifecycle triage → todo → ready → running → blocked → done → archived. The research child should enter running first (it runs on the researcher profile's model), then hand off so the script task becomes ready, and so on down the pipeline. A task sitting in blocked is asking for a human — open it and read why.

9. Review each task's workspace deliverables Click into each completed child task. Every task gets its own workspace (a scratch directory / dedicated dir / git worktree, depending on configuration) where its outputs live. Open the researcher task's workspace and read the brief; open the scriptwriter's and read the script + scene list; open the video_producer's and play the rendered video. Confirm each stage's output actually fed the next — the script should quote the brief's key points, and the video should follow the scene list.

10. Accept the final video to done If the video meets the brief, accept the top-level task so it moves to done (and later archived). If it doesn't, add a comment on the relevant child task with concrete feedback (e.g. "narration too fast in scene 2") and send that task back through the board rather than redoing everything — that is the point of decomposed work.

11. Reflect: why this beats one big agent Look back at the run: the expensive long-context model only did research, the cheap model wrote the script, and the tool model rendered. One prompt to a single agent would have used the expensive model for everything. Note one sentence on cost and one on reviewability (you could inspect and reject per stage).

12. The natural-language way You can skip the CLI entirely and simply ask the agent: "Create a crew of profiles — researcher, scriptwriter, video producer — then decompose this board task and route each piece to the right profile".

Test it

Three profiles exist, the decomposed Kanban tasks route to the right profile by description, and the accepted board delivers a rendered video.

Note: Full commands and screenshots are in labs/lab13-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 14 — Webhook

Learning outcome: LO3: Integrate the agent with external systems via webhooks..

Goal: Wire the agent to the outside world with webhooks: trigger the agent from an incoming webhook payload and/or have it call an outgoing webhook, then secure the endpoint.

What you'll build

A webhook that triggers the agent (or that the agent calls), verified end-to-end. (Tools: Hermes Agent, webhooks.)

Step-by-step

1. Confirm the local gateway is serving Webhooks ride on the local gateway — if it isn't up, nothing can receive a request. Check its status and note the host/port it reports (typically localhost plus a port); that host/port is the base of your webhook URL. Or just ask the agent: "Check the gateway is up so webhooks can reach you"

`hermes gateway status`

2. Create / enable a webhook endpoint on the dashboard Open the dashboard and click Webhooks (/webhooks) in the sidebar. Click New webhook (or the enable toggle), give it a recognizable name like briefing-trigger, and save. The dashboard generates a unique endpoint URL for it. > The exact page layout and creation flow are version-specific — verify in the video / Hermes docs (<https://hermes-agent.nousresearch.com/docs/>).

3. Copy the endpoint URL and map it to an agent action Use the copy button next to the new webhook to copy its full URL, and paste it somewhere handy. Then set what the webhook does: map incoming payloads to an agent action — for this lab, have it run a short briefing (e.g. instruction: "When this webhook fires, read the payload's message field and act on it"). Without a mapping, a payload is accepted but nothing happens.

4. Send a test payload with curl From a terminal, simulate an external system by POSTing a JSON payload to the endpoint. Replace <endpoint-url> with the URL you copied in step 4: Or just ask the agent: "POST a test payload to the webhook endpoint and confirm a 2xx response"

```
curl -X POST <endpoint-url> \  
  
-H "Content-Type: application/json" \  
  
-d '{"event":"test","message":"Trigger Athena briefing"}'
```

5. Watch the agent react Switch to the chat/session view (or the logs) and confirm the payload actually triggered Athena: the mapped action runs and its output appears — e.g. a briefing message referencing "Trigger Athena briefing" from your payload. Also check the webhook's row on the /webhooks page; most builds show a delivery/last-triggered indicator. Accepted-but-no-action means the mapping in step 4 isn't wired — fix it and re-send.

6. Secure the endpoint, then re-send with the secret An open webhook lets anyone on the network drive your agent, so lock it down. On the webhook's settings, add a shared secret/token (and an IP allowlist if your build offers one), save, and re-send the test including the secret header — it should still return 2xx and trigger the agent: Or just ask the agent: "Re-send the test payload with the Bearer secret header and confirm it's accepted"

```
curl -X POST <endpoint-url> \  
  
-H "Authorization: Bearer <secret>" \  
  
-H "Content-Type: application/json" \  
  
-d '{"event":"test","message":"Trigger Athena briefing"}'
```

7. Confirm unauthorized calls are rejected Re-run the step 5 curl (the one without the secret header). It must now be rejected with 401/403, and Athena must not act. If it still returns 2xx, the secret isn't enforced — re-check the webhook's security settings and save again. You now have a working, secured inbound integration.

Test it

A test payload to the webhook triggers the agent's mapped action, and calls without the secret are rejected.

Note: Full commands and screenshots are in labs/lab14-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Topic 02 — OpenClaw — Automating a Business Back-Office (33%)

Install · Providers · Channels · Skills · Tools · Commands · Crons · Memory · Dreaming · Security · Multi-Agent · Use Cases

Key concepts

- OpenClaw runs on Node.js 24 LTS as a long-running gateway daemon that hosts every channel, tool, cron and skill for the Nimbus Supplies back office.
- Providers are connected by OAuth sign-in (OpenAI Codex, MiniMax) or API key (Anthropic, DeepSeek, OpenRouter); skills from skills.sh and integrations such as Firecrawl and AgentMail extend the agent.
- 'Dreaming' lets the agent reflect during idle time — reviewing memory and generating follow-ups — while profiles and allow/deny lists keep every channel least-privileged and audited.
- Crons schedule recurring work, a heartbeat auto-restarts the gateway, and a multi-agent capstone splits work across cooperating Sales, Research and Ops agents for real back-office use cases.

Lab 15 — Install OpenClaw

Learning outcome: LO1: Install and configure an autonomous AI agent platform locally and on a container/VPS..

Goal: The learner installs the OpenClaw platform on their own machine (or a cloud VPS) by first installing Node.js 24 LTS, then installing OpenClaw via npm and running the onboarding wizard, and verifies the install with doctor and gateway status. This gateway becomes the back-office brain of the Nimbus Supplies running example.

What you'll build

A running OpenClaw gateway daemon (local or VPS) verified with openclaw doctor (Tools: OpenClaw, Node.js 24 LTS, npm, openclaw gateway.)

Step-by-step

1. Install Node.js 24 LTS (24 recommended; 22.16+ minimum). OpenClaw runs on Node. Windows — download the Node 24 LTS installer from <<https://nodejs.org/>> and run the .msi (WSL2 is recommended for stability), then check:

```
node -v
```

```
npm -v
```

2. Install OpenClaw — Option A (npm, recommended, cross-platform). After Node is installed:

```
npm install -g openclaw@latest
```

```
openclaw onboard
```

3. Install OpenClaw — Option B (installer script, fallback if npm fails). macOS / Linux / WSL2:

```
curl -fsSL https://openclaw.ai/install.sh | bash
```

4. Install OpenClaw — Option C (from source, advanced). Requires git and pnpm:

```
git clone https://github.com/openclaw/openclaw.git
```

```
cd openclaw
```

```
pnpm install && pnpm build && pnpm ui:build
```

```
pnpm link --global
```

```
openclaw onboard --install-daemon
```

5. (Optional) Install on a VPS so Nimbus Supplies runs 24/7. The same three-step flow works on any Ubuntu host — a Hostinger KVM VPS or an exe.dev sandbox.

SSH in first (ssh root@<your-vps-ip>), then run:

```
# Step 1 – Node.js 24
```

```
curl -fsSL https://deb.nodesource.com/setup_24.x | sudo -E bash -
```

```
sudo apt-get install -y nodejs
```

```
node -v && npm -v
```

```
# Step 2 – OpenClaw (note: sudo for a global install on a fresh VPS)
```

```
sudo npm install -g openclaw@latest
```

```
# Step 3 – onboarding wizard
```

```
openclaw onboard
```

6. Verify the installation.

```
openclaw --version
```

```
openclaw doctor
```

```
openclaw gateway status
```

7. Confirm the gateway is running (and set to auto-start). The gateway is the long-running daemon that will host every channel, tool, and cron for Nimbus Supplies:

```
openclaw gateway status
```

```
openclaw gateway start # if it is not already running
```

Test it

openclaw --version prints a version, openclaw doctor shows green checks for Node 24 and the gateway daemon (provider/channel checks pending until Labs 16-17), and openclaw gateway status reports running.

Note: Full commands and screenshots are in labs/lab15-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 16 — Models & Providers

Learning outcome: LO2: Connect an LLM model/provider to an agent via OAuth or API key..

Goal: The learner connects a language model to OpenClaw, trying the supported provider options — OAuth sign-ins (OpenAI Codex, MiniMax) and API-key providers (Anthropic, DeepSeek, OpenRouter) — learns where credentials are stored under ~/.openclaw/auth/, hot-swaps between models, and runs a Nimbus Supplies smoke test.

What you'll build

A Nimbus Supplies agent wired to a working model provider that passes openclaw model test (Tools: OpenClaw, OpenAI Codex / MiniMax (OAuth), Anthropic / DeepSeek / OpenRouter (API key).)

Step-by-step

1. See what is available and what is currently selected. Or just ask the agent: "Which models can you use, and which one is active right now?"

```
openclaw models list
```

```
openclaw model current
```

2. Option A — OpenAI Codex (OAuth, GPT-5.5). OAuth uses the openai-codex provider so you sign in with your ChatGPT Plus / Pro subscription instead of paying per token. Credentials are stored in ~/.openclaw/auth/ (profile at ~/.openclaw/auth-profiles/openai-codex.json); refresh is automatic.

```
openclaw models auth login --provider openai-codex
```

3. Option B — MiniMax (OAuth, MiniMax-M2.7). Enable the OAuth plugin, restart the gateway, then log in and set it as default:

```
openclaw plugins enable minimax-portal-auth
```

```
openclaw gateway restart
```

```
openclaw models auth login --provider minimax-portal --set-default
```

4. Option C — Anthropic (API key, Claude Opus 4.7). Get a key from <<https://console.anthropic.com/settings/keys>>: Or just ask the agent: "Switch your model to Anthropic's Claude Opus 4.7" Or just ask the agent: "Say hello so I can confirm your model connection works"

```
export ANTHROPIC_API_KEY="sk-ant-..."

openclaw config set provider anthropic

openclaw config set model anthropic/claude-opus-4-7

openclaw model test
```

5. Option D — DeepSeek (API key, V4). Get a key from <https://platform.deepseek.com/api_keys>:

```
export DEEPSEEK_API_KEY="sk-..."

openclaw config set provider deepseek

openclaw config set model deepseek/deepseek-v4

openclaw model test
```

6. Option E — OpenRouter (API key, routes to Claude Opus 4.7). Get a key from <<https://openrouter.ai/keys>>:

```
export OPENROUTER_API_KEY="sk-or-..."

openclaw config set provider openrouter

openclaw config set model openrouter/anthropic/claude-opus-4.7

openclaw model test
```

7. Persist API keys so they survive reboots. For the API-key providers, add the export line to your shell profile (~/.zshrc on macOS, ~/.bashrc on Linux). OAuth tokens are persisted automatically in ~/.openclaw/auth/ and do not need this.

```
echo 'export ANTHROPIC_API_KEY="sk-ant-..."' >> ~/.zshrc

source ~/.zshrc
```

8. Switch models and confirm the active one. You can hot-swap the global default at any time: Or just ask the agent: "Make DeepSeek v4 your default model from now on"

```
openclaw model use deepseek/deepseek-v4

openclaw model current

openclaw model test
```

9. Give the agent a Nimbus Supplies smoke test. With any model active, from the CLI chat (openclaw) ask: > You are the back-office assistant for Nimbus Supplies, a

small office-supplies reseller. In two sentences, introduce yourself to a customer. A sensible reply confirms the model is wired up. Or just ask the agent: "Introduce yourself in two sentences as the Nimbus Supplies back-office assistant"

Test it

openclaw models list shows the connected models, openclaw model current prints the active one, openclaw model test returns a short successful completion with no auth error, and OAuth credentials appear under ~/.openclaw/auth/.

Note: Full commands and screenshots are in labs/lab16-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 17 — Channels

Learning outcome: LO2: Connect multiple messaging channels to one agent gateway..

Goal: The learner gives customers a way to reach the Nimbus Supplies agent by creating a Telegram bot end-to-end via BotFather, registering it with a token, pairing a WhatsApp number by scanning a QR code, and running both channels at once through a single OpenClaw gateway so the same agent answers on either platform.

What you'll build

Telegram and WhatsApp both live on one OpenClaw gateway, answered by the same Nimbus Supplies agent (Tools: OpenClaw, Telegram (BotFather), WhatsApp, openclaw channel.)

Step-by-step

1. Create a Telegram bot with BotFather. 1. Open Telegram and search for @BotFather. 2. Send /newbot. 3. Enter a display name (e.g. Nimbus Supplies Assistant). 4. Enter a username — must be unique and end in _bot (e.g. nimbus_supplies_bot). 5. BotFather replies with an HTTP API token like 123456789:ABC-DEF1234ghIkl-zyx57W2v1u123ew11. Copy it.
2. Register and start the Telegram channel. Or just ask the agent: "Connect Telegram using the bot token I got from BotFather" Or just ask the agent: "Start the Telegram channel and confirm it is live"

```
openclaw channel add telegram --token 123456789:ABC-DEF1234ghIkl-zyx57W2v1u123ew11
```

```
openclaw channel start telegram
```

3. Test Telegram. In Telegram, search for your bot's username and send: > Hi, do you supply recycled A4 paper for Nimbus Supplies? The agent should reply using the model you connected in Lab 16.

4. Add and start the WhatsApp channel. Or just ask the agent: "Add WhatsApp as a channel and show me the pairing QR code" Or just ask the agent: "Start the WhatsApp channel so I can scan the QR code"

```
openclaw channel add whatsapp
```

```
openclaw channel start whatsapp
```

5. Pair WhatsApp by scanning the QR. On your phone: WhatsApp → Settings → Linked Devices → Link a Device, then scan the QR code shown in the terminal. Wait for whatsapp: connected in the OpenClaw log. > WhatsApp keeps a stateful session on disk at ~/.openclaw/whatsapp/. Do not delete that folder unless you intend to re-pair.

6. Confirm both channels are running at the same time. One gateway hosts many channels — the same Nimbus Supplies agent now answers on both. Or just ask the agent: "Which channels are you listening on right now?"

```
openclaw channel list
```

```
openclaw gateway status
```

7. Inspect and manage a channel. Useful day-to-day commands:

```
openclaw channel show telegram
```

```
openclaw channel restart whatsapp # e.g. after the QR expires
```

Test it

openclaw channel list shows telegram and whatsapp both running, a message to the Telegram bot gets an agent reply, a message to the linked WhatsApp account gets an agent reply, and openclaw gateway status shows both connected.

Note: Full commands and screenshots are in labs/lab17-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 18 — Skills

Learning outcome: LO3: Extend an agent with registry and custom skills..

Goal: The learner installs and removes skills from the skills.sh registry, invokes them from chat, and builds a custom Nimbus Supplies skill (nimbus-quote) under ~/.openclaw/skills/ so the agent handles a recurring back-office task — drafting an itemised customer quote with GST — the same way every time.

What you'll build

Registry skills plus a custom nimbus-quote skill that produces itemised, GST-inclusive quotes (Tools: OpenClaw, skills.sh registry, ~/.openclaw/skills/, openclaw skills.)

Step-by-step

1. List the skills currently installed. Or just ask the agent: "What skills do you have installed?"

```
openclaw skills list
```

2. Install a few useful skills from the registry. Skills come from <<https://skills.sh/>>. Or just ask the agent: "Install the web-research skill from the skills.sh registry"

```
openclaw skills add web-research --source skills.sh
```

```
openclaw skills add self-improvement --source skills.sh
```

3. Invoke a skill from chat. In your Telegram or WhatsApp channel, send:

```
/skill web-research "Who are the main office-supplies wholesalers in Singapore?"
```

4. Remove a skill you do not need (keeps the agent focused and safe):

```
openclaw skills remove self-improvement
```

5. Refresh the registry index if a skill you expect is missing:

```
openclaw skills update
```

6. Build a custom Nimbus Supplies skill. Skills live under `~/.openclaw/skills/`. Create a folder and a skill definition. (File layout is shown here as a documented example — confirm the exact schema for your version at <<https://docs.openclaw.ai/>> before relying on advanced fields.) Or just ask the agent: "Create a new custom skill called nimbus-quote for drafting itemised customer quotes"

```
mkdir -p ~/.openclaw/skills/nimbus-quote
```

7. Register and test the custom skill. Or just ask the agent: "Reload your skills so the new nimbus-quote skill is picked up"

```
openclaw skills list # confirm nimbus-quote appears
```

```
openclaw gateway restart # reload skills if it is not picked up
```

8. The natural-language way You can skip the CLI entirely and simply ask the agent: "Install the github skill from ClawHub and show me what it can do".

Test it

`openclaw skills list` shows the registry skills plus the custom nimbus-quote, `/skill web-research` runs end-to-end, and `/skill nimbus-quote` produces an itemised quote with a 9% GST line and standard terms, marking any unknown price as TBC instead of guessing.

Note: Full commands and screenshots are in `labs/lab18-*/README.md`. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 19 — Tools & Integrations

Learning outcome: LO3: Extend an agent with third-party tool integrations..

Goal: The learner extends the Nimbus Supplies agent with two real integrations — Firecrawl (JS-rendered web scraping) and AgentMail (its own inbox/outbox) — then runs an end-to-end back-office task: research a supplier's website prices, draft an itemised quote with the Lab 18 skill, and email it to a customer. Tool profiles and allow/deny lists scope tools per channel.

What you'll build

A Firecrawl + AgentMail integrated agent that scrapes supplier prices and emails an itemised quote (Tools: OpenClaw, Firecrawl, AgentMail, openclaw tools.)

Step-by-step

1. Inspect the built-in tools first. Or just ask the agent: "What tools do you have, grouped by category?"

```
openclaw tools list
```

2. Enable Firecrawl. Sign up at <<https://www.firecrawl.dev/>>, then Dashboard → API Keys → copy the fc-... key: Or just ask the agent: "Enable the Firecrawl scraping tool with my API key"

```
export FIRECRAWL_API_KEY="fc-..."
```

```
openclaw tools enable firecrawl --api-key "$FIRECRAWL_API_KEY"
```

```
openclaw tools status firecrawl
```

3. Smoke-test Firecrawl from chat. In Telegram or WhatsApp: > Use Firecrawl to fetch the products page of a stationery supplier's website and list their A4 paper options with prices in 5 bullets.
4. Enable AgentMail. Create an inbox at <<https://agentmail.to/>>, copy the API key and inbox address: Or just ask the agent: "Set up AgentMail with my API key so you have your own inbox"

```
export AGENTMAIL_API_KEY="..."
```

```
openclaw tools enable agentmail \
```

```
--api-key "$AGENTMAIL_API_KEY" \
```

```
--inbox bot@yourname.agentmail.to
```

```
openclaw tools status agentmail
```

5. Smoke-test AgentMail. From chat: > Send an email from my AgentMail inbox to <your-personal-email> with subject "Hello from Nimbus Supplies" and a friendly one-

paragraph body. Check your inbox, reply to it, then ask: > Check my AgentMail inbox and summarise the latest reply.

6. Run the real end-to-end task (research → quote → email). From a channel, send one instruction that chains both tools plus the skill from Lab 18: > A customer, Acme Cafe, wants 10 reams of recycled A4 paper and 5 boxes of black pens. Use Firecrawl to check current prices on our supplier's site, then use the nimbus-quote skill to draft an itemised quote, and email it from my AgentMail inbox to orders@acmecafe.example with subject "Your Nimbus Supplies quote".

7. Scope tools per channel with profiles and allow/deny lists. Public channels should not run shell exec. Apply the safe messaging preset and lock down a public channel: Or just ask the agent: "On the Telegram channel, block shell command execution" Or just ask the agent: "On Telegram, allow only the web tool group"

```
openclaw channel set telegram --profile messaging
```

```
openclaw channel set telegram --deny exec
```

```
openclaw channel set telegram --allow group:web
```

```
openclaw channel show telegram
```

8. (Optional) Add one more integration of your choice. Pick one from <<https://docs.openclaw.ai/tools>> and enable it (exact name/flags per the docs), e.g.:

```
openclaw tools enable <tool-name> --api-key <KEY>
```

9. The natural-language way You can skip the CLI entirely and simply ask the agent: "Use the browser to pull today's top three AI headlines and save them to headlines.md".

Test it

openclaw tools list --enabled shows firecrawl and agentmail, the agent scrapes a supplier URL and returns structured prices, an AgentMail send arrives in a real inbox, the end-to-end task produces an itemised quote email, and a channel set to --deny exec refuses shell-exec requests.

Note: Full commands and screenshots are in labs/lab19-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 20 — Commands

Learning outcome: LO4: Operate an agent through its CLI and in-chat command surface..

Goal: The learner becomes fluent in the two ways to drive OpenClaw daily: the CLI commands (config, model, doctor, gateway, channel, tools, skills, cron) and the in-chat slash commands operators and customers use inside a channel, building a personal

quick-reference table for running the Nimbus Supplies back office and streaming live gateway logs.

What you'll build

A personal quick-reference of the CLI and slash commands used to run Nimbus Supplies day to day (Tools: OpenClaw, openclaw CLI, in-chat slash commands.)

Step-by-step

1. `openclaw config` — read and change settings. Or just ask the agent: "Show me your current configuration settings"

```
openclaw config list                # show all settings

openclaw config get model           # get one setting

openclaw config set model anthropic/claude-opus-4-7

openclaw config edit                # open the config file in
$EDITOR
```

2. `openclaw model` — manage the brain.

```
openclaw model list

openclaw model current

openclaw model use deepseek/deepseek-v4

openclaw model test
```

3. `openclaw doctor` — health check (with auto-fix).

```
openclaw doctor

openclaw doctor --fix
```

4. `openclaw gateway` — control the daemon that runs everything.

```
openclaw gateway status

openclaw gateway start

openclaw gateway stop

openclaw gateway restart

openclaw gateway logs --tail 50

openclaw gateway install    # install daemon (LaunchAgent / systemd /
Task Scheduler)
```

5. `openclaw channel` — manage the front doors.

```
openclaw channel list
```

```
openclaw channel add telegram --token <TOKEN>
```

```
openclaw channel start whatsapp
```

```
openclaw channel show telegram
```

```
openclaw channel restart whatsapp
```

```
openclaw channel remove telegram
```

6. openclaw tools and openclaw skills — capabilities.

```
openclaw tools list
```

```
openclaw tools list --enabled
```

```
openclaw tools status agentmail
```

```
openclaw skills list
```

```
openclaw skills add web-research --source skills.sh
```

```
openclaw skills remove web-research
```

7. In-chat slash commands. These work inside any channel (Telegram, WhatsApp, CLI chat). Try each from your Nimbus Supplies bot:

8. Watch it work live. In one terminal, follow the logs; in the channel, send a message and observe the corresponding lines:

```
openclaw gateway logs --follow
```

Test it

You can switch the model both from chat (/model, per-conversation) and globally from the CLI (openclaw model use), /whoami returns your platform user ID, openclaw gateway logs --follow streams live activity, and /help lists the available slash commands.

Note: Full commands and screenshots are in labs/lab20-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 21 — Cron Jobs & Heartbeat

Learning outcome: LO4: Automate agent work with scheduled crons and a monitoring heartbeat..

Goal: The learner makes Nimbus Supplies run unattended by scheduling recurring tasks with cron (a nightly sales report and a weekly supplier price-check using Firecrawl) and monitoring uptime with the heartbeat, so the gateway self-checks, auto-restarts on failure, and can page an external monitor if it dies.

What you'll build

A nightly sales-report cron, a weekly price-check cron, and an enabled self-healing heartbeat (Tools: OpenClaw, openclaw cron, gateway heartbeat, Firecrawl + AgentMail.)

Step-by-step

1. Create the nightly sales-report cron. Cron syntax is <min> <hour> <day> <month> <weekday>. Or just ask the agent: "Every night at 9pm, summarise today's customer chats and quotes and post it to Telegram"

```
# Every day at 21:00 – post a Nimbus Supplies end-of-day summary to Telegram
```

```
openclaw cron create \  
  --schedule "0 21 * * *" \  
  --prompt "Summarise today's Nimbus Supplies customer chats and any quotes sent. List open questions that still need a human." \  
  --channel telegram \  
  --name nightly-sales-report
```

2. Create a weekly supplier price-check cron (uses Firecrawl from Lab 19):

```
# Every Monday at 08:00 – check supplier prices and email me
```

```
openclaw cron create \  
  --schedule "0 8 * * 1" \  
  --prompt "Use Firecrawl to fetch our main supplier's A4 paper price and email me the current price via AgentMail with subject 'Weekly supplier price'." \  
  --channel telegram \  
  --name weekly-price-check
```

3. List and inspect your crons. Or just ask the agent: "What cron jobs do you have scheduled?"

```
openclaw cron list
```

```
openclaw cron show nightly-sales-report
```

```
openclaw cron logs nightly-sales-report --tail 20
```

4. Fire a cron on demand to test it now (ignores the schedule): Or just ask the agent: "Run the nightly sales report right now so I can check it works"

```
openclaw cron run nightly-sales-report
```

5. Disable / enable / delete crons as your needs change:

```
openclaw cron disable weekly-price-check
```

```
openclaw cron enable weekly-price-check
```

```
openclaw cron delete weekly-price-check # only if you no longer need it
```

6. Enable the heartbeat. The heartbeat is a periodic self-check that verifies the gateway is alive, the model responds, and channels are connected; on failure the daemon attempts auto-restart and logs the incident. Or just ask the agent: "Turn on your heartbeat self-check every 60 seconds" Or just ask the agent: "Is your heartbeat healthy? Give me the status"

```
openclaw gateway heartbeat enable --interval 60s
```

```
openclaw gateway heartbeat status
```

```
openclaw gateway heartbeat logs --tail 20
```

7. (Optional) Push a heartbeat ping to an external monitor so you are paged if Nimbus Supplies goes dark (e.g. a free <<https://healthchecks.io/>> check):

```
openclaw gateway heartbeat webhook \
```

```
--url https://hc-ping.com/<your-uuid> \
```

```
--interval 5m
```

8. Confirm auto-restart works. Ensure the daemon is installed to auto-start, then simulate a crash: Or just ask the agent: "Shut down your gateway - I want to check the heartbeat restarts you"

```
openclaw gateway status # expect "running" + "auto-start: yes"
```

```
openclaw gateway stop
```

```
# wait ~60s for the service manager to restart it
```

```
openclaw gateway status # expect "running" again
```

9. The natural-language way You can skip the CLI entirely and simply ask the agent: "Add a heartbeat task: every 30 minutes, check my inbox and flag anything urgent".

Test it

openclaw cron list shows nightly-sales-report and weekly-price-check, openclaw cron run delivers a summary to Telegram immediately, openclaw gateway heartbeat status reports healthy, and after openclaw gateway stop the service manager restarts the gateway within about a minute.

Note: Full commands and screenshots are in labs/lab21-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 22 — Memory & Context

Learning outcome: LO2: Give an agent persistent, user-scoped memory..

Goal: The learner gives the Nimbus Supplies agent durable memory so it recalls customer preferences across conversations: using the in-chat `/memory` command, inspecting where OpenClaw persists state under `~/.openclaw/`, teaching a customer's standing preference, clearing the conversation to prove it is real memory, and confirming recall in a later separate chat.

What you'll build

An agent that persistently remembers a customer's standing preferences and applies them in a fresh conversation (Tools: OpenClaw, `/memory`, `/clear`, `~/.openclaw/` state.)

Step-by-step

1. See what OpenClaw stores on disk. Or just ask the agent: "Show me what you store in your `~/.openclaw` state folder"

```
ls -la ~/.openclaw/
```

2. Inspect your current memory from chat. In your Telegram or WhatsApp bot:

```
/memory
```

3. Teach the agent a durable customer preference. Still in chat, send a plain-language instruction: > Remember this for future chats: our customer Acme Cafe always wants recycled A4 paper (never standard), delivery on Tuesdays only, and quotes addressed to Mei from Acme. Then confirm it was captured:

```
/memory
```

4. Clear the conversation to prove it is real memory, not just context.

```
/clear
```

5. Recall in a fresh conversation. In the now-cleared chat, ask something that requires the remembered facts: > Draft a quote for Acme Cafe: 10 reams of A4 paper and 5 boxes of black pens. A correct agent should automatically choose recycled A4, note Tuesday delivery, and address the quote to Mei — without you repeating any of it. (This pairs perfectly with the nimbus-quote skill from Lab 18.)
6. Edit or correct a memory. If a preference changes, update it in natural language and re-check: > Update your memory: Acme Cafe now accepts delivery on Tuesdays or Thursdays.

```
/memory
```

7. Back up memory before major changes (memory lives under ~/.openclaw/, so the Lab 15 backup habit applies): Or just ask the agent: "Back up all your memory and state to a dated folder in my home directory"

```
mkdir -p ~/openclaw-backup-$(date +%F)

cp -R ~/.openclaw/ ~/openclaw-backup-$(date +%F)/openclaw-state
2>/dev/null || true
```

Test it

/memory shows the customer preferences after you teach them, after /clear the /memory list still shows them (short-term context cleared, long-term memory retained), and a fresh quote request for the customer automatically applies recycled paper, Tuesday delivery and the correct recipient name without repetition.

Note: Full commands and screenshots are in labs/lab22-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 23 — Security

Learning outcome: LO4: Secure an agent with key management, allowlists, sandboxing and audit..

Goal: The learner hardens the Nimbus Supplies deployment before real customers: storing API keys as environment variables (never in cleartext config), restricting who can DM the bot with allowlists, applying per-channel tool deny lists, tightening the code-execution sandbox (timeout/memory/network), reviewing the audit log, and rotating provider keys.

What you'll build

A hardened agent with keys in env, DM allowlists, per-channel deny lists, a capped sandbox and an exported audit log (Tools: OpenClaw, allowlists, tool deny lists, exec/python sandbox, gateway logs.)

Step-by-step

1. Secure API keys — never hardcode or commit them. Store keys once as environment variables. macOS / Linux — append to ~/.zshrc or ~/.bashrc:

```
export ANTHROPIC_API_KEY="sk-ant-..."

export FIRECRAWL_API_KEY="fc-..."

export AGENTMAIL_API_KEY="..."
```

2. Restrict who can DM the bot (allowlists). By default anyone who finds your Telegram bot can message it. Lock it to known users: Or just ask the agent: "Only

allow Telegram users 12345678 and 87654321 to message you; block everyone else"

```
openclaw channel set telegram \  
    --allow-users 12345678,87654321 \  
    --pair-mode allowlist
```

3. Apply tool deny lists (least privilege per channel). A customer-facing channel has no business running shell commands or deleting files: Or just ask the agent: "On Telegram, deny shell exec and any file writes or deletes"

```
openclaw channel set telegram --deny exec,fs.write,fs.delete  
  
openclaw channel set telegram --allow group:web  
  
openclaw channel set telegram --profile messaging  
  
openclaw channel show telegram | grep -E "(allow|deny|profile)"
```

4. Tighten the code-execution sandbox. The exec / Python tools run sandboxed — cap time, memory, and network: Or just ask the agent: "Cap your exec sandbox at 10 seconds, 256MB memory and no network access"

```
openclaw tools config exec \  
    --timeout 10s \  
    --memory 256mb \  
    --network deny  
  
openclaw tools config python \  
    --timeout 30s \  
    --memory 512mb \  
    --network deny
```

5. Review the audit log. Every tool call, model call, and channel event is logged:

```
openclaw gateway logs --tail 100  
  
openclaw gateway logs --filter tool=exec  
  
openclaw gateway logs --since 1h --filter level=warn  
  
openclaw gateway logs --export ~/openclaw-audit-$(date +%F).log
```

6. Rotate provider keys on a schedule (e.g. every 90 days). 1. Generate a new key in the provider dashboard. 2. Update the environment variable. 3. Restart the

gateway: openclaw gateway restart. 4. Confirm openclaw model test still works. 5. Revoke the old key in the provider dashboard.

Test it

cat ~/.openclaw/config.toml does not contain raw API keys, a user not on the Telegram allowlist gets a not-authorized response, openclaw channel show telegram lists the deny list and messaging profile, and the audit log shows blocked exec attempts on a locked channel.

Note: Full commands and screenshots are in labs/lab23-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 24 — Multi-Agent Workforce

Learning outcome: LO5: Orchestrate cooperating multi-agent teams to deliver a business outcome..

Goal: In the OpenClaw capstone the learner splits the single agent into three cooperating agents — Sales (customer-facing on Telegram/WhatsApp), Research (supplier scouting with Firecrawl), and Ops (email, quotes, nightly crons) — each with least-privilege tools, and runs one end-to-end business scenario across all three, with a verified fallback using isolated OPENCLAW_HOME instances that cooperate via AgentMail and a shared folder.

What you'll build

A three-agent Sales/Research/Ops workforce that runs a customer order end-to-end into a real quote email (Tools: OpenClaw, openclaw agent (or isolated OPENCLAW_HOME instances), Firecrawl, AgentMail.)

Step-by-step

1. Create three named agents. Attempt the native multi-agent path first (confirm exact syntax in the docs): Or just ask the agent: "Create a new agent named Sales to handle customer-facing chats"

```
openclaw agent create sales
```

```
openclaw agent create research
```

```
openclaw agent create ops
```

```
openclaw agent list
```

2. Give each agent least-privilege tools (mirrors Lab 19/23 allow-deny, applied per agent): Or just ask the agent: "Give the Sales agent web and quoting tools only, with no shell or file writes" Or just ask the agent: "Restrict the Research agent to Firecrawl, web search and browser only"

```
# Sales – web + quoting only, no shell, no filesystem writes
```

```
    openclaw agent set sales --allow group:web --deny  
exec,fs.write,fs.delete
```

```
# Research – scraping only
```

```
    openclaw agent set research --allow firecrawl,web_search,browser --deny  
exec,fs.write
```

```
# Ops – email + files, no public channel exposure
```

```
openclaw agent set ops --allow agentmail,group:fs
```

3. Attach channels to the right agents. Or just ask the agent: "Route the Telegram channel to the Sales agent"

```
# Customer traffic goes to Sales
```

```
openclaw channel set telegram --agent sales
```

```
openclaw channel set whatsapp --agent sales
```

```
# A private research bot for the operator (see Lab 17 for creating a 2nd  
bot)
```

```
openclaw channel add telegram --token <RESEARCH_BOT_TOKEN>
```

```
openclaw channel set <research-bot> --agent research
```

4. Give each agent the right skill/memory. Sales uses nimbus-quote and the customer memory from Lab 22; Research uses nimbus-supplier-brief (Lab 18 exercise). Install/enable as needed:

```
openclaw skills list
```

5. Wire cooperation between agents. The simplest, robust hand-off uses the tools you already have:

- Sales → Ops: when Sales captures an order, it asks Ops (via a shared note file or an AgentMail message) to produce and send the quote.
- Ops → Research: when a price is unknown ("TBC"), Ops asks Research to look it up with Firecrawl.
- Research → Ops: Research replies with the price; Ops finalises the quote email. Or just ask the agent: "Every night at 9pm, compile today's orders and quotes and post them to Telegram"

```
# Ops nightly consolidation cron (from Lab 21, now owned by Ops)
```

```
openclaw cron create \
```

```
--schedule "0 21 * * *" \

--prompt "Compile today's Nimbus Supplies orders and quotes sent, and
list anything waiting on Research." \

--channel telegram \

--name ops-nightly-report
```

6. Fallback (fully verified) — run agents as separate profiles/instances. If native multi-agent commands are not in your build, model each agent as its own OpenClaw configuration using only Lab 15–19 primitives: - Sales = your main instance with Telegram/WhatsApp, --profile messaging, deny exec,fs.write (Lab 23). - Research = a second instance pointed at a separate config dir with only Firecrawl/web tools enabled:

```
OPENCLAW_HOME=~/.openclaw-research openclaw onboard
```

```
OPENCLAW_HOME=~/.openclaw-research openclaw tools enable firecrawl --
api-key "$FIRECRAWL_API_KEY"
```

7. Run the end-to-end capstone scenario. From a customer's Telegram, trigger the whole workflow: > Hi Nimbus Supplies — I'm Mei from Acme Cafe. I need 10 reams of recycled A4 paper and 5 boxes of black pens. What's your best price and when can you deliver? Expected chain: 1. Sales greets Mei, recalls Acme's preferences from memory (recycled paper, Tuesday delivery — Lab 22), and captures the order. 2. Sales hands the order to Ops, which starts the nimbus-quote. The pen price is unknown → marked TBC. 3. Ops asks Research to fetch the current pen price via Firecrawl. 4. Research replies with the price; Ops finalises the itemised quote (with GST) and emails it to Acme via AgentMail. 5. That night, the ops-nightly-report cron summarises the day, including Mei's order.

8. The natural-language way You can skip the CLI entirely and simply ask the agent: "Spawn two sub-agents to research our top two competitors in parallel, then compare their pricing".

Test it

Three agents (or isolated instances) each have a distinct least-privilege tool set (Sales cannot exec, Research cannot write files, Ops is not on a public channel), a customer Telegram message produces a real quote email combining memory, the nimbus-quote skill, a Firecrawl price lookup and AgentMail delivery, and the nightly cron posts a summary referencing the same order.

Note: Full commands and screenshots are in labs/lab24-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 25 — Dreaming — Idle Reflection & Self-Improvement

Learning outcome: LO3: Use OpenClaw 'dreaming' so the agent reflects during idle time..

Goal: Enable OpenClaw's 'dreaming' feature so the Nimbus Supplies agent uses idle time to review its memory, consolidate what it has learned, and generate useful follow-up actions and skill improvements — then inspect what it produced. Reference the OpenClaw functions playlist for the walkthrough.

What you'll build

An agent that, when idle, 'dreams' — reflecting on memory and proposing follow-ups you can review. (Tools: OpenClaw, dreaming, memory.)

Step-by-step

1. Enable the dreaming feature in OpenClaw config. This flips the setting that lets the gateway schedule reflection cycles when the agent is idle: Or just ask the agent: "Turn on dreaming so you reflect and consolidate memory when idle"

```
openclaw config set dreaming.enabled true
```

2. Let the agent sit idle so a dream cycle runs — or trigger one manually. Rather than wait for a natural idle window, force a cycle now so you can observe it immediately: Or just ask the agent: "Run a dream cycle now instead of waiting for idle time"

```
openclaw dream run
```

3. Review what the dream produced. List the insights, follow-ups, and skill suggestions the cycle generated: Or just ask the agent: "Show me the insights and follow-ups from your last dream"

```
openclaw dream list
```

4. Accept a useful follow-up the agent proposed. In the dashboard (or via the CLI shown in the video), open one proposal that genuinely helps Nimbus Supplies and approve it. Accepting promotes the proposal into a real action or a saved skill; declining discards it. This human review step keeps you in control of what the agent's downtime actually changes.

Test it

A dream cycle runs during idle time and produces reviewable insights/follow-ups from the agent's memory.

Note: Full commands and screenshots are in labs/lab25-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 26 — Use Cases & Key Functions

Learning outcome: LO3: Apply OpenClaw's key functions to real back-office use cases..

Goal: Tie the OpenClaw track together with practical use cases for Nimbus Supplies — a Google Workspace CLI assistant, a data-analytics agent, and a social-media-marketing agent — combining channels, skills, tools, memory, dreaming and crons into end-to-end automations. Each use case has its own reference video.

What you'll build

One or more end-to-end OpenClaw automations (Workspace, analytics or social media) solving a real use case. (Tools: OpenClaw, skills, tools, Google Workspace, crons, memory.)

Step-by-step

1. Watch one of the use-case videos above and study how the presenter composes several functions into one automation — which channel triggers the flow, which skills and tools it calls, and how memory and crons fit in. Good starting points: a Google Workspace CLI assistant (Docs/Sheets/Gmail), a Data Analytics agent, or a Social Media Marketing agent.
2. Pick a back-office use case. Choose one concrete Nimbus Supplies workflow to implement end-to-end. A strong default is "research a supplier, then email a quote": a customer asks for a price, the agent looks up current supplier pricing, applies a margin, and emails a formatted quote. Other options: customer FAQ Q&A on Telegram, or an automated nightly orders-and-quotes report.
3. Compose the channels, skills and tools needed for it. Wire the pieces together for your chosen use case. For research-to-quote that means: the customer channel (e.g. Telegram) as the trigger, the skill that formats a quote (e.g. nimbus-quote from Lab 18), a research tool (Firecrawl / web_search) for live pricing, and a delivery tool (AgentMail) to send the email. Confirm the agent has permission to call each one.
4. Run the automation end-to-end and verify the outcome. Send a real trigger — e.g. message the agent "What's your price for 50 kg of arabica beans?" on the customer channel — and watch the full chain fire: the agent recalls context from memory, runs the price lookup tool, builds the quote with the skill, and delivers it via email. Confirm the actual deliverable (the sent quote email) is correct, not just that the logs look busy.
5. Schedule it or hand it to the multi-agent team from Lab 24. Make the automation ongoing. Either attach it to a cron (e.g. a nightly consolidation report) so it runs unattended, or delegate the use case to the Sales / Research / Ops agents you built in Lab 24 so the right specialist owns each step. This is what turns a one-off demo into part of an always-on Nimbus Supplies back office.

Test it

A real use case (Google Workspace, data analytics or social media) runs end-to-end using OpenClaw's channels, skills, tools and crons.

Note: Full commands and screenshots are in labs/lab26-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Topic 03 — Paperclip — Running & Governing a Company of AI Agents (34%)

Install (Windows & Mac) · Company, CEO & Mission · Adaptors · Task Backlog · Tavily Search API · Hire the Team · Daily Routine

Key concepts

- Paperclip is an operating system for a company of AI agents, self-hosted with Docker Compose on Windows (WSL2) or macOS; you are the Board, a CEO agent reports to you, and specialists report to the CEO.
- Running use case: 'Alterra AI Blogs' — a zero-human company whose mission is to research the AI landscape and publish AI-related blogs end to end; the mission you write steers every agent it hires.
- Model adaptors (built-in Claude Code, Codex and Gemini CLI) are the agents' engine; the Tavily Search API — stored as a Secret and bound to a runtime env var — gives researchers live web results.
- The company runs off a detailed task backlog (todo -> in_progress -> in_review -> done) whose core task hires the agents to publish the blogs end to end — every hire pauses at a Board approval gate — and a daily 3pm routine generates and runs the publish task automatically.

Lab 27 — Install Paperclip on Windows & Mac

Learning outcome: LO1: Install and self-host Paperclip on Windows and macOS..

Goal: Watch the Paperclip overview, then self-host Paperclip with Docker Compose on your own machine — Docker Desktop on macOS, or Docker Desktop + WSL2 on Windows — and open the dashboard at <http://localhost:3100>.

What you'll build

A running Paperclip instance at <http://localhost:3100> on Windows or Mac. (Tools: Paperclip, Docker Desktop, Docker Compose.)

Step-by-step

1. Install Docker Desktop — macOS: download from docker.com; Windows: enable WSL2 first, then install Docker Desktop with the WSL2 backend - macOS —

download Docker Desktop from [docker.com](https://www.docker.com/products/docker-desktop/) (<https://www.docker.com/products/docker-desktop/>), drag it into Applications, launch it and wait for the whale icon to settle. - Windows — enable WSL2 first (open PowerShell as Administrator), then install Docker Desktop and select the WSL2 backend during setup:

```
wsl --install
```

2. Clone the Paperclip repository

```
git clone https://github.com/paperclipai/paperclip.git
```

3. Start Paperclip with the quickstart compose file (same command in the macOS Terminal or the Windows WSL2/PowerShell shell)

```
docker compose -f docker-compose.quickstart.yml up --build
```

4. Open the dashboard in your browser at <http://localhost:3100> and create your Board account Browse to <http://localhost:3100> (<http://localhost:3100>). The first-run screen asks you to create your Board account — this identity is the human owner every agent ultimately answers to. Sign in and land on the empty dashboard.

Test it

Docker shows the Paperclip containers running and the dashboard loads at <http://localhost:3100> on your platform.

Note: Full commands and screenshots are in [labs/lab27-*/README.md](#). Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 28 — Setup Company, CEO & Mission

Learning outcome: LO2: Found the company — name, mission and the CEO agent..

Goal: Found 'Altera AI Blogs': create the company, write its mission (research the AI landscape daily and publish AI-related blogs end to end), set the monthly budget, then hire the CEO agent — the Chief Officer who will run the blog operation and report to you, the Board.

What you'll build

A company with a mission and budget, plus an approved CEO agent at the top of the org chart. (Tools: Paperclip, company settings, org chart.)

Step-by-step

1. Create the company 'Altera AI Blogs' from the company switcher In the dashboard open the company switcher (top of the sidebar) and choose New company. Name it exactly: > Altera AI Blogs This is the zero-human company you will govern for the rest of the topic.

2. Write the mission — it steers every agent the company hires In the company setup (or Company Settings → Mission), write the mission statement: > Research the AI landscape daily and publish AI-related blogs end to end The mission is not decoration: every agent the company hires reads it and aligns its work to it, so a vague mission produces a vague company.
3. Set the monthly budget cap in Company Settings Open Company Settings and set the monthly budget cap (e.g. \$50/month to start). The budget is the hard ceiling on what the whole company may spend on model calls — agents cannot exceed it, which makes it your first governance lever as the Board.
4. Hire the CEO agent (Chief Officer) and approve the hire — you are the Board Hire the company's first agent: the CEO (Chief Officer). The hire pauses at an approval gate — as the Board, review the CEO's mandate and click Approve. The CEO is the only agent that reports directly to you; every later hire will report to the CEO.
5. Open the Org page and confirm the structure: Board above the CEO Open the Org page and check the org chart: Board (you) at the top, the CEO directly beneath. This two-node chart is the seed of the whole workforce you will grow in Labs 30–33.

Test it

The company exists with a mission and budget, and an approved CEO agent sits under the Board on the org chart.

Note: Full commands and screenshots are in labs/lab28-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 29 — Setup Adaptor

Learning outcome: LO2: Connect a model adaptor as the agents' engine..

Goal: Give every Altera AI Blogs agent a brain. An adaptor is the bridge between Paperclip and the language model that does the reasoning: when any agent (the CEO or a specialist) picks up a task, Paperclip hands the work to the adaptor, which runs it on the underlying model and streams the result back to the task. Paperclip ships three BUILT-IN adaptors that piggyback on locally installed CLIs — Claude Code (claude_local, 9 models), Codex (codex_local, 10 models) and Gemini CLI (gemini_local, 8 models) — and supports installable EXTERNAL adaptor packages (alpha). You install/log in to the CLI on the host so Paperclip can detect it, confirm it shows as available under Settings → Instance settings → Adapters, use the power icon to hide adaptors you don't want agents to pick, and restart Paperclip if detection fails. Whatever adaptor you enable here powers every hire in Labs 30-33.

What you'll build

A detected, enabled model adaptor (Claude Code, Codex or Gemini CLI) powering every Altera AI Blogs agent. (Tools: Paperclip, model adaptors (Claude Code / Codex / Gemini CLI).)

Step-by-step

1. Install / log in to the adaptor CLI on the host so Paperclip can detect it Paperclip detects adaptor CLIs installed on the host machine. Install and authenticate at least one — e.g. Claude Code (claude and log in) or the OpenAI Codex CLI (codex and log in) — on the same machine that runs the Paperclip containers.
2. Open Settings → Instance settings → Adapters in the dashboard In the dashboard go to Settings → Instance settings → Adapters. This page is the registry of every engine the instance can run agents on.
3. Confirm the built-in adaptors are listed and yours shows as available You should see the three built-ins with their model counts: - Claude Code (claude_local) — 9 models - Codex (codex_local) — 10 models - Gemini CLI (gemini_local) — 8 models The adaptor whose CLI you installed in step 2 should show as available/detected.
4. Use the power icon to hide adaptors; 'Install Adapter' adds external adaptor packages (alpha) Two controls on this page matter for governance: - The power icon on each adaptor toggles it — hide adaptors you don't want agents to pick, so hires can only run on engines you have vetted. - Install Adapter adds external adaptor packages beyond the three built-ins. This feature is alpha — expect rough edges.
5. Restart Paperclip if your adaptor is not detected Detection happens at startup, so a CLI installed after Paperclip came up may not show until a restart:

```
docker compose -f docker-compose.quickstart.yml restart
```

Test it

Settings → Adapters lists the built-in adaptors and your chosen adaptor is detected and enabled for agents.

Note: Full commands and screenshots are in labs/lab29-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 30 — Create the Task Backlog

Learning outcome: LO3: Create a detailed task backlog that drives the company..

Goal: Write the backlog that runs Altera AI Blogs. Every task gets a precise title AND a detailed description — the description is the agent's brief, so vague tasks produce vague work. The CORE task is 'Hire the agents to publish the AI-related blogs end to end'; the rest build the blog pipeline around it.

What you'll build

A Tasks board holding six well-specified tasks, with the end-to-end blog-publishing hire as the core task. (Tools: Paperclip, Tasks board.)

Step-by-step

1. Create the CORE task — 'Hire the agents to publish the AI-related blogs end to end' — and assign it to the CEO Title: > Hire the agents to publish the AI-related blogs end to end Description: > Based on the content strategy, propose which agents to hire (research analyst, blog writer, editor/publisher, SEO/marketer) with roles, budgets and reporting lines, so the company can research, write, edit and publish AI blogs without human writers. Assign it to the CEO. This is the core task of the whole topic — in Lab 32 the CEO works it to staff the company.
2. Create 'Define the AI blog content strategy and editorial plan' Title: > Define the AI blog content strategy and editorial plan Description: > Pick the beats (models, chips, policy, funding), the publishing cadence, the blog format and the quality bar for every published post.
3. Create 'Scaffold the blog workspace & tooling foundation' Title: > Scaffold the blog workspace & tooling foundation Description: > Set up the workspace folder structure, post templates, style guide and source register the whole team will use.
4. Create 'Build the AI news sources watchlist' Title: > Build the AI news sources watchlist Description: > Assemble the RSS feeds, publication sites, X accounts, arXiv categories and company blogs to monitor, with priority tiers and de-duplication rules.
5. Create 'Produce and publish the first AI blog post' Title: > Produce and publish the first AI blog post Description: > Pull the top stories from the watchlist, verify each against two independent sources, draft a cited post and publish it for Board review.
6. Create 'Wire the research pipeline to live search' Title: > Wire the research pipeline to live search Description: > Integrate the Tavily Search API so researchers query the live web with fresh results (depends on the Tavily key from the next lab).

Test it

The Tasks board shows all six tasks, each with a detailed description, and the core hiring task is assigned to the CEO.

Note: Full commands and screenshots are in labs/lab30-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 31 — Add the Tavily Search API

Learning outcome: LO2: Wire live web search into the company via the Tavily Search API..

Goal: A blog company needs fresh information. Create a Tavily API key, store it as a company Secret, and bind it to the TAVILY_API_KEY runtime environment variable —

Paperclip resolves the secret server-side when a run starts, so the key never sits in a prompt or a repo.

What you'll build

A TAVILY_API_KEY secret bound to the agents' runtime environment, verified with a live search task. (Tools: Paperclip, Secrets, Tavily Search API.)

Step-by-step

1. Create a Tavily account and copy an API key from app.tavily.com Sign up at app.tavily.com (<https://app.tavily.com>), open the dashboard and copy an API key (it starts with tvly-). Keep it in your clipboard — you will paste it once, into a secret, and nowhere else.
2. Open Settings → Company settings → Secrets and click 'New secret' In the Paperclip dashboard go to Settings → Company settings → Secrets and click New secret. Secrets are company-scoped: every agent the company hires can be granted them, but no agent ever reads them in plain text.
3. Store the key as a secret (e.g. name it tavily) Name the secret (e.g. tavily), paste the API key as its value, and save. Secrets are injected at run start, never shown to agents in plain text — the value will not appear in prompts, task logs or the repo.
4. Bind it to TAVILY_API_KEY via an agent's Environment variables (or a project's Env field) Now map the secret to the environment variable the Tavily SDK/CLI expects: - Open an agent's Environment variables (or a project's Env field), - add the key TAVILY_API_KEY, - choose Secret as the value type, and - select the stored tavily secret. At run start Paperclip resolves the binding server-side and the agent's runtime sees TAVILY_API_KEY set — without the key ever passing through a prompt.
5. Verify with a live search task Create a task for an agent with the binding: > Use Tavily to find today's three biggest AI announcements and list your sources The agent should return fresh, cited results — stories from today with source links, not stale training-data knowledge.

Test it

The secret exists, TAVILY_API_KEY is bound to the runtime env, and a live-search task returns fresh cited results.

Note: Full commands and screenshots are in `labs/lab31-*/README.md`. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 32 — Hire the Members Under the Hiring Task

Learning outcome: LO3: Staff the company through the hiring task, behind approval gates..

Goal: Let Altera AI Blogs staff itself — all of it happening UNDER the core task from Lab 30, 'Hire the agents to publish the AI-related blogs end to end'. You move that task to in_progress and the CEO works it: inside the task it drafts a hiring plan proposing each specialist the blog pipeline needs — a research analyst (scans the watchlist and verifies stories), a blog writer (drafts cited posts), an editor/publisher (reviews, formats and publishes) and an SEO/marketer (titles, tags and distribution) — each with a mandate, a budget cap and a reporting line to the CEO. Every proposed hire pauses at an approval gate for you, the Board: approve it and the agent is created; reject it and the CEO revises the plan. When the gates clear, the new members appear in the Agents sidebar and on the Org chart beneath the CEO, and you assign them the rest of the backlog so the pipeline runs end to end — research → draft → edit → publish.

What you'll build

A Board-approved blog team (research analyst, blog writer, editor/publisher, SEO/marketer) hired under the core task, visible on the org chart and taking the backlog. (Tools: Paperclip, hiring task, approval gates, org chart.)

Step-by-step

1. Open the CORE task 'Hire the agents to publish the AI-related blogs end to end' and move it to in_progress On the Tasks board open the CORE task 'Hire the agents to publish the AI-related blogs end to end' and move it to in_progress. The CEO starts working the brief and drafts the hiring plan inside this task — watch the task's activity log as the plan takes shape.
2. Review the proposed roles in the task's hiring plan — research analyst, blog writer, editor/publisher, SEO/marketer Read the CEO's plan as a Board member would. It should propose the blog-pipeline specialists — research analyst, blog writer, editor/publisher, SEO/marketer — and for each: the mandate, a per-agent budget, and the reporting line (everyone reports to the CEO).
3. Approve each proposed hire at its approval gate (or reject and ask for a revised proposal) — you are the Board Each proposed hire pauses at an approval gate. For every one, either: - Approve — the agent is hired and comes online, or - Reject with a comment — the CEO revises the proposal (e.g. lower budget, sharper mandate) and resubmits. Nothing joins the company without an explicit Board decision — this gate is the human-in-the-loop control at the heart of Paperclip governance.
4. Confirm the new members appear under Agents and on the Org chart beneath the CEO Open Agents and check every approved specialist is listed, then open the Org chart and confirm the structure: Board → CEO → the four specialists beneath the CEO.
5. Assign the rest of the backlog to the new team and watch the pipeline run: research → draft → edit → publish Assign the remaining Lab 30 tasks to the right specialists — the content strategy and watchlist to the research analyst, the workspace scaffold and the first post to the blog writer, review and publication to the

editor/publisher, titles/tags/distribution to the SEO/marketer — and watch the board: the pipeline runs research → draft → edit → publish, with tasks flowing todo → in_progress → in_review → done, the CEO coordinating and you approving anything that hits a gate.

Test it

The core hiring task reaches in_review/done, every hire was Board-approved under that task, and the specialists appear on the org chart publishing the first posts.

Note: Full commands and screenshots are in labs/lab32-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Lab 33 — Routine & Trigger — Publish Daily at 3pm

Learning outcome: LO3: Automate the company with a scheduled routine and trigger..

Goal: Put Altera AI Blogs on autopilot. Create a Routine with a daily 3:00pm trigger that generates and runs the blog-publishing task automatically — the routine fires, creates the day's 'Produce and publish today's AI blog post' task, the team executes the pipeline (research → draft → edit → publish), and the finished post lands in in_review for the Board. No human kicks it off.

What you'll build

A daily 3:00pm routine that generates and runs the blog-publishing task automatically. (Tools: Paperclip, Routines, triggers, Tasks board.)

Step-by-step

1. Open the Routines page from the sidebar (Work → Routines) In the dashboard sidebar open Work → Routines. This page holds every scheduled automation the company runs — it is the equivalent of a cron board for your agent workforce.
2. Create a new routine: 'Daily AI blog publish' Click New routine and name it: > Daily AI blog publish
3. Set the trigger: schedule, daily at 3:00pm (15:00) In the routine's trigger section choose schedule and set it to daily at 3:00pm (15:00). The trigger is what fires the routine — from now on the clock, not a human, starts the day's publishing run.
4. Define what the routine generates: the task 'Produce and publish today's AI blog post' Configure the routine to generate the task: > Produce and publish today's AI blog post Description: > Pull the top stories from the watchlist, verify against two sources, draft a cited post, publish for Board review — assigned to the team. Every time the trigger fires, a fresh copy of this task lands on the board and the team executes it end to end.
5. Save and enable the routine, then trigger it once manually to confirm it creates and runs the task Save the routine and switch it to enabled, then use the manual

run/trigger now control to fire it once. Confirm on the Tasks board that the task was created and the team started the pipeline: research → draft → edit → publish.

6. At the next 3:00pm, confirm the task was generated and executed automatically. Wait for the next 3:00pm tick and check: the task 'Produce and publish today's AI blog post' was generated and executed with no human action, and the finished post is sitting in in_review for you, the Board.

Test it

The routine is enabled with a daily 15:00 trigger, a manual run creates and executes the publish task, and the next 3pm run happens without human action.

Note: Full commands and screenshots are in labs/lab33-*/README.md. Use only accounts, keys and hosts you own, and keep agents under human oversight.

Wrap-Up — The Agent Build Arc and Governance

These cross-cutting themes run through every platform in the course. Study this section alongside the labs so the same skillset transfers from Hermes to OpenClaw to Paperclip and into your own production agents.

The agent build arc

Every platform follows the same progression — the order in which you built each agent in the labs:

- Install — set up the runtime (Hermes CLI/TUI, the OpenClaw Node.js gateway daemon, or Paperclip via Docker Compose).
- Models — connect an LLM provider (Claude, OpenAI, OpenRouter, MiniMax or DeepSeek) and pick a capable model.
- Channels — reach the agent where work happens (Telegram, WhatsApp, Discord, Slack and more).
- Memory — give the agent durable, user- or company-scoped recall across sessions and surfaces.
- Skills & tools — extend the agent with reusable skills and third-party tool/integration calls.
- Crons & heartbeat — schedule recurring work and let a heartbeat keep the agent running.
- Security — isolate execution in sandboxes, store secrets in config, and gate risky actions with approvals.
- Multi-agent — orchestrate a coordinator plus specialist workers to deliver an end-to-end business outcome.

Human-in-the-loop governance

- Approvals — important or destructive actions pause for explicit human sign-off before they run.
- Budgets — company- and per-agent spend caps warn at 80% and pause at 100% so costs never run away.
- Sandboxing — agents execute in isolated backends (Docker, SSH, or a remote host) so the host stays safe.
- Audit — an activity log and cost-events trail record who did what, giving you a CFO-style view of the workforce.
- Least privilege — scope each channel, skill and tool with allow/deny lists so agents can only do their job.

Next Steps

- First pass: complete every lab on your own machine, following the commands in each lab file.
- Second pass: redo the labs from memory until installing, connecting and securing an agent is automatic.
- Extend the labs with your own skills, tools and integrations for a workflow you care about.
- Deploy an agent on a VPS or cloud backend for 24/7 operation, with budgets, approvals and audit turned on.
- Explore multi-agent orchestration — a coordinator plus specialists — for a realistic end-to-end business task.
- Review each lab's detailed steps in this guide and re-run the labs on your own machine.

Glossary

- **Autonomous AI agent** — An LLM-driven system that perceives, reasons, acts with tools and remembers, running under human oversight.
- **Agent stack** — The layers that make an agent work: model, memory, skills, tools, channels and a scheduler/security layer.
- **Model / provider** — The LLM (and the vendor serving it) that gives the agent its reasoning — e.g. Claude, OpenAI, OpenRouter.
- **Channel** — A messaging surface the agent is reachable on, such as Telegram, WhatsApp, Discord or Slack.
- **Memory** — Durable, user- or company-scoped storage that lets the agent recall context across sessions and surfaces.
- **Skill** — A reusable, named capability an agent can install or create and invoke on demand.

- **Tool / integration** — An external action or service the agent can call — web search, image generation, email, GitHub via MCP.
- **MCP** — Model Context Protocol — a standard way to connect an agent to external tool servers.
- **Cron / heartbeat** — A schedule that triggers recurring agent work; a heartbeat self-checks and keeps the agent running.
- **Sandbox** — An isolated execution backend (Docker, SSH, remote host) that contains what an agent's commands can touch.
- **Approval gate** — A human sign-off required before an agent performs an important or destructive action.
- **Budget cap** — A spend limit at company or per-agent level that warns at 80% and pauses work at 100%.
- **Multi-agent orchestration** — A coordinator agent delegating to specialist worker agents to deliver an end-to-end outcome.